

Parker Christenson Assignment 2.1: Text Classification using Word Embeddings

Instructions

Prior to beginning work on this assignment, review the week's lab session.

Dataset

The dataset comprises newsgroup documents categorized into 20 different topics. Each document is labeled with its corresponding newsgroup category.

Download the dataset from the following link:

<https://www.kaggle.com/datasets/crawford/20-newsgroups?resource=download>

- Load the "20 Newsgroups" dataset into a pandas DataFrame.
- Preprocess the text data: remove stopwords, punctuation, convert to lowercase, etc.
- Split the data into training and testing sets.
- Utilize pre-trained word embeddings (e.g., Word2Vec, GloVe) and/or train your own embeddings using the training data.
- Build a text classification model using a classifier of your choice (e.g., Logistic Regression, Support Vector Machine, Neural Network).
- Train the model using the transformed training data (with embeddings).
- Predict the categories for the testing data.
- Evaluate the model's performance
- Evaluate the model's performance using the following classification metrics:
 - Accuracy
 - Precision (weighted)
 - Recall (weighted)
 - F1 Score (weighted)
 - Confusion Matrix
 - Area Under the Receiver Operating Characteristic curve (AUC-ROC)

```
In [100.. # text pre processing imports
import re
import string
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.stem import PorterStemmer
from nltk.corpus import wordnet
from nltk import ne_chunk, pos_tag
from nltk.tree import Tree
from nltk import pos_tag
```

```
from tqdm import tqdm
import os

# standard imports
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# sklearn imports
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import roc_auc_score

# tensorflow imports
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.utils import to_categorical
```

```
In [49]: print("Tensorflow version: ", tf.__version__)
print("The GPU is", "available" if tf.config.experimental.list_physical_de
print("Num GPUs Available: ", len(tf.config.experimental.list_physical_de
```

```
Tensorflow version: 2.17.0
The GPU is available
Num GPUs Available: 1
```

```
In [50]: # Loading the data set from sklearn
from sklearn.datasets import fetch_20newsgroups
newsgroups = fetch_20newsgroups(subset='all')

print("There are total of", len(newsgroups.target_names), "categories")
```

```
There are total of 20 categories
```

```
In [51]: # displaying some of the data in the dataset
print(newsgroups.data[0])
```

From: Mamatha Devineni Ratnam <mr47+@andrew.cmu.edu>
 Subject: Pens fans reactions
 Organization: Post Office, Carnegie Mellon, Pittsburgh, PA
 Lines: 12
 NNTP-Posting-Host: po4.andrew.cmu.edu

I am sure some bashers of Pens fans are pretty confused about the lack of any kind of posts about the recent Pens massacre of the Devils. Actually, I am a bit puzzled too and a bit relieved. However, I am going to put an end to non-Pittsburghers' relief with a bit of praise for the Pens. Man, they are killing those Devils worse than I thought. Jagr just showed you why he is much better than his regular season stats. He is also a lot of fun to watch in the playoffs. Bowman should let JAGR have a lot of fun in the next couple of games since the Pens are going to beat the pulp out of Jersey anyway. I was very disappointed not to see the Islanders lose the final regular season game. PENS RULE!!!

```
In [52]: # printing the target names
print(newsgroups.target_names)
```

```
['alt.atheism', 'comp.graphics', 'comp.os.ms-windows.misc', 'comp.sys.ibm.pc.hardware', 'comp.sys.mac.hardware', 'comp.windows.x', 'misc.forsale', 'rec.autos', 'rec.motorcycles', 'rec.sport.baseball', 'rec.sport.hockey', 'sci.crypt', 'sci.electronics', 'sci.med', 'sci.space', 'soc.religion.christian', 'talk.politics.guns', 'talk.politics.mideast', 'talk.politics.misc', 'talk.religion.misc']
```

```
In [53]: # getting a count of the target names and their frequency
unique, counts = np.unique(newsgroups.target, return_counts=True)
print(np.asarray((unique, counts)).T)
```

```
[[ 0 799]
 [ 1 973]
 [ 2 985]
 [ 3 982]
 [ 4 963]
 [ 5 988]
 [ 6 975]
 [ 7 990]
 [ 8 996]
 [ 9 994]
 [10 999]
 [11 991]
 [12 984]
 [13 990]
 [14 987]
 [15 997]
 [16 910]
 [17 940]
 [18 775]
 [19 628]]
```

```
In [54]: # making the data into a dataframe
data = pd.DataFrame({'text': newsgroups.data, 'target': newsgroups.target})
```

```
print(data.head())
```

	text	target
0	From: Mamatha Devineni Ratnam <mr47+@andrew.cm...	10
1	From: mblawson@midway.ecn.uoknor.edu (Matthew ...	3
2	From: hilmi-er@dsv.su.se (Hilmi Eren)\nSubject...	17
3	From: guyd@austin.ibm.com (Guy Dawson)\nSubjec...	3
4	From: Alexander Samuel McDiarmid <am2o+@andrew...	4

```
In [55]: # printing the first text in the data
print(data['text'][5])
```

```
From: tell@cs.unc.edu (Stephen Tell)
Subject: Re: subliminal message flashing on TV
Organization: The University of North Carolina at Chapel Hill
Lines: 25
NNTP-Posting-Host: rukbat.cs.unc.edu
```

In article <7480237@hpfco.FC.HP.COM> myers@hpfco.FC.HP.COM (Bob Myers) writes:

```
>> Hi. I was doing research on subliminal suggestion for a psychology
>> paper, and I read that one researcher flashed hidden messages on the
>> TV screen at 1/200ths of a second. Is that possible?
```

```
> Might
> even be a vector ("strokewriter") display, in which case the lower limit
> on image time is anyone's guess (and is probably phosphor-persistence lim
ited).
```

Back in high school I worked as a lab assistant for a bunch of experimental psychologists at Bell Labs. When they were doing visual perception and memory experiments, they used vector-type displays, with 1-millisecond refresh rates common.

So your case of 1/200th sec is quite practical, and the experimenters were probably sure that it was 5 milliseconds, not 4 or 6 either.

```
>Bob Myers KC0EW >myers@fc.hp.com
```

Steve

--

```
Steve Tell      tell@cs.unc.edu H: 919 968 1792 | #5L Estes Park apts
UNC Chapel Hill Computer Science W: 919 962 1845 | Carrboro NC 27510
Engineering is a _lot_ like art: Some circuits are like lyric poems, some
are like army manuals, and some are like The Hitchhiker's Guide to the Gal
axy..
```

```
In [70]: # downloading the punkt_tab
nltk.download('punkt')
nltk.download('wordnet')
nltk.download('punkt_tab')
nltk.download('averaged_perceptron_tagger_eng')
nltk.download('maxent_ne_chunker_tab')
nltk.download('words')
```

```

[nltk_data] Downloading package punkt to /home/parker/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /home/parker/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package punkt_tab to /home/parker/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger_eng to
[nltk_data] /home/parker/nltk_data...
[nltk_data] Package averaged_perceptron_tagger_eng is already up-to-
[nltk_data] date!
[nltk_data] Downloading package maxent_ne_chunker_tab to
[nltk_data] /home/parker/nltk_data...
[nltk_data] Package maxent_ne_chunker_tab is already up-to-date!
[nltk_data] Downloading package words to /home/parker/nltk_data...
[nltk_data] Unzipping corpora/words.zip.

```

Out[70]: True

```

In [71]: def Preprocess(text_data):
        """
        Function to preprocess text file, generate preprocessed email domain,
        """

        preprocessed_email = []
        preprocessed_subject = []
        preprocessed_text = []

        for sentence in tqdm(text_data):

            # Preprocessing email
            domain = re.findall("@[\w.]+", sentence)
            email = ""
            for items in domain:
                items = items.replace("@", "")
                items = items.split(".")
                for i in set(items):
                    if((len(i) > 2) and i != "com" and i != "COM"):
                        email += i + " "
            preprocessed_email.append(email.strip())

            # Preprocessing subject
            subject = "" # Initialize subject as an empty string
            text_split = sentence.split("\n")
            for item in text_split:
                if item.startswith("Subject:"):
                    for word in item.split():
                        if not word.endswith(":"):
                            subject += word + " "
                    subject = re.sub("[^0-9a-zA-Z\s]", " ", subject)
                    subject = " ".join(subject.split()).strip()
            preprocessed_subject.append(subject.lower())

            # Preprocessing text
            text = re.sub(r"(.*)Subject:(.*?)(.*)\n", " ", sentence) # Remo
            text = re.sub(r"(.*)From:(.*?)(.*)\n", " ", text) # Remo
            text = re.sub(r"(.*)Write to:(.*?)(.*)\n", " ", text) # Remo
            text = re.sub(r"(.*):(.*?)", " ", text) # Remo

            # Decontraction
            text = re.sub(r"won't", "will not", text)

```

```

text = re.sub(r"can\t", "can not", text)
text = re.sub(r"n\t", " not", text)
text = re.sub(r"\re", " are", text)
text = re.sub(r"\s", " is", text)
text = re.sub(r"\d", " would", text)
text = re.sub(r"\ll", " will", text)
text = re.sub(r"\t", " not", text)
text = re.sub(r"\ve", " have", text)
text = re.sub(r"\m", " am", text)

text = re.sub(r"[\w\-.]+\.[\w\-.]+\b", " ", text) # Remo
text = re.sub(r"[\n\t]", " ", text) # Remo
text = re.sub(r"<.*>", " ", text) # Remo
text = re.sub(r"\(.*\)", " ", text) # Remo
text = re.sub(r"\[.*\]", " ", text) # Remo
text = re.sub(r"\{.*\}", " ", text) # Remo
text = re.sub(r"[0-9]+", " ", text) # Remo
text = re.sub(r"^[A-Za-z\s]", " ", text) # Remo

# Named Entity Recognition and Replacement
chunks = list(ne_chunk(pos_tag(word_tokenize(text))))

for i in chunks:
    if isinstance(i, Tree):
        if i.label() == "GPE":
            j = i.leaves()
            if len(j) > 1: # If a city name has two or more wor
                gpe = "_".join([term for term, pos in j])
                text = re.sub(rf"{j[1][0]}", gpe, text)
                text = re.sub(rf"{j[0][0]}", " ", text)
            if i.label() == "PERSON":
                for term, pos in i.leaves():
                    text = re.sub(re.escape(term), "", text)

# Clean up underscores
text = re.sub(r"\b([a-zA-Z]+)\b", r"\1", text) # Replace _word
text = re.sub(r"\b([a-zA-Z]+)\b", r"\1", text) # Replace _word
text = re.sub(r"\b([a-zA-Z]+)\b", r"\1", text) # Replace word_
text = re.sub(r"\b[a-zA-Z]{1}_([a-zA-Z]+)", r"\1", text)
text = re.sub(r"\b[a-zA-Z]{2}_([a-zA-Z]+)", r"\1", text)

text = text.lower()
text_split = text.split()
text = " ".join(word for word in text_split if 2 < len(word) < 15)
preprocessed_text.append(text.strip())

return preprocessed_email, preprocessed_subject, preprocessed_text

```

```

In [72]: # applying the function to the text column
emails, subjects, texts = Preprocess(data['text'])

```

```

100%|██████████| 18846/18846 [1:08:49<00:00, 4.56it/s]

```

```

In [77]: # making the final dataframe with the subjects and texts columns
final_data = pd.DataFrame({'email': emails, 'subject': subjects, 'text':
final_data.head()

```

Out[77]:

	email	subject	text	target
0	edu cmu andrew	pens fans reactions	office carnegie mellon pittsburgh andrew cmu e...	10
1	edu midway uoknor ecn edu essex uoknor ecn	which high performance vlb video card	seek for video card midway ecn uoknor edu engi...	3
2	dsv dsv	armenia says it could shoot down turkish plane...	viktorias dsv dept computer and would stretch f...	17
3	austin ibm austin ibm pal500 julian uwo heart ...	ide vs scsi dma and detach	ibm can anyone explain fairly simple terms why...	3
4	edu cmu andrew	driver	sophomore mechanical engineering carnegie mell...	4

Building the text classification model using tensorflow

```
In [79]: # using word2vec to convert the text into vectors after tokenizing
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
```

```
In [81]: # tokenize the text data
sentences = [word_tokenize(text) for text in data['text']]

# train the Word2Vec model
word2vec_model = Word2Vec(sentences, vector_size=100, window=5, min_count

# printing the number of words in the model
print(f'Trained Word2Vec model with {len(word2vec_model.wv)} words in the
```

Trained Word2Vec model with 276455 words in the vocabulary.

```
In [82]: # function to conver the text
def text_to_embedding(text, model, embedding_dim=100):
    words = word_tokenize(text)
    word_vectors = [model.wv[word] for word in words if word in model.wv]

    if not word_vectors: # if no word retun none
        return np.zeros(embedding_dim)

    return np.mean(word_vectors, axis=0)

# apply the function to the text col
data['text_embedding'] = data['text'].apply(lambda x: text_to_embedding(x
```

```
In [85]: X = np.array(data['text_embedding'].tolist())
y = final_data['target']

# splitting the data for the testing and training split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

```
In [93]: # setting the number of topics
num_topics = len(set(y_train))

# building the model
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(X_train.shape[1],)))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_topics, activation='softmax'))

# Compile the model
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', m
```

```
/home/parker/anaconda3/envs/tf_ml/lib/python3.10/site-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [94]: # Train the model
model.fit(X_train, y_train, epochs=100, batch_size=64, validation_split=0
```


Epoch 1/100
189/189 - 6s - 29ms/step - accuracy: 0.0993 - loss: 2.8669 - val_accuracy: 0.1983 - val_loss: 2.5268
Epoch 2/100
189/189 - 0s - 1ms/step - accuracy: 0.1912 - loss: 2.4187 - val_accuracy: 0.2974 - val_loss: 2.0984
Epoch 3/100
189/189 - 0s - 1ms/step - accuracy: 0.2402 - loss: 2.2201 - val_accuracy: 0.3253 - val_loss: 2.0165
Epoch 4/100
189/189 - 0s - 1ms/step - accuracy: 0.2734 - loss: 2.1200 - val_accuracy: 0.3501 - val_loss: 1.9224
Epoch 5/100
189/189 - 0s - 1ms/step - accuracy: 0.2980 - loss: 2.0479 - val_accuracy: 0.3528 - val_loss: 1.8736
Epoch 6/100
189/189 - 0s - 1ms/step - accuracy: 0.3185 - loss: 1.9996 - val_accuracy: 0.3840 - val_loss: 1.8023
Epoch 7/100
189/189 - 0s - 1ms/step - accuracy: 0.3345 - loss: 1.9431 - val_accuracy: 0.4208 - val_loss: 1.7514
Epoch 8/100
189/189 - 0s - 1ms/step - accuracy: 0.3514 - loss: 1.9133 - val_accuracy: 0.4171 - val_loss: 1.7295
Epoch 9/100
189/189 - 0s - 1ms/step - accuracy: 0.3547 - loss: 1.8831 - val_accuracy: 0.4188 - val_loss: 1.7254
Epoch 10/100
189/189 - 0s - 1ms/step - accuracy: 0.3621 - loss: 1.8670 - val_accuracy: 0.4377 - val_loss: 1.6781
Epoch 11/100
189/189 - 0s - 1ms/step - accuracy: 0.3710 - loss: 1.8374 - val_accuracy: 0.4380 - val_loss: 1.6658
Epoch 12/100
189/189 - 0s - 1ms/step - accuracy: 0.3749 - loss: 1.8288 - val_accuracy: 0.4254 - val_loss: 1.6717
Epoch 13/100
189/189 - 0s - 1ms/step - accuracy: 0.3750 - loss: 1.8152 - val_accuracy: 0.4486 - val_loss: 1.6544
Epoch 14/100
189/189 - 0s - 1ms/step - accuracy: 0.3883 - loss: 1.7900 - val_accuracy: 0.4446 - val_loss: 1.6438
Epoch 15/100
189/189 - 0s - 1ms/step - accuracy: 0.3905 - loss: 1.7903 - val_accuracy: 0.4463 - val_loss: 1.6387
Epoch 16/100
189/189 - 0s - 1ms/step - accuracy: 0.3971 - loss: 1.7686 - val_accuracy: 0.4277 - val_loss: 1.6792
Epoch 17/100
189/189 - 0s - 1ms/step - accuracy: 0.4001 - loss: 1.7607 - val_accuracy: 0.4632 - val_loss: 1.6067
Epoch 18/100
189/189 - 0s - 1ms/step - accuracy: 0.4046 - loss: 1.7328 - val_accuracy: 0.4622 - val_loss: 1.5923
Epoch 19/100
189/189 - 0s - 1ms/step - accuracy: 0.4020 - loss: 1.7475 - val_accuracy: 0.4542 - val_loss: 1.6007
Epoch 20/100
189/189 - 0s - 1ms/step - accuracy: 0.4108 - loss: 1.7368 - val_accuracy: 0.4493 - val_loss: 1.5830

Epoch 21/100
189/189 - 0s - 1ms/step - accuracy: 0.4038 - loss: 1.7266 - val_accuracy: 0.4649 - val_loss: 1.5814
Epoch 22/100
189/189 - 0s - 1ms/step - accuracy: 0.4154 - loss: 1.7205 - val_accuracy: 0.4721 - val_loss: 1.5718
Epoch 23/100
189/189 - 0s - 1ms/step - accuracy: 0.4164 - loss: 1.7143 - val_accuracy: 0.4569 - val_loss: 1.5873
Epoch 24/100
189/189 - 0s - 1ms/step - accuracy: 0.4175 - loss: 1.6962 - val_accuracy: 0.4755 - val_loss: 1.5709
Epoch 25/100
189/189 - 0s - 1ms/step - accuracy: 0.4241 - loss: 1.6950 - val_accuracy: 0.4589 - val_loss: 1.5723
Epoch 26/100
189/189 - 0s - 1ms/step - accuracy: 0.4227 - loss: 1.6906 - val_accuracy: 0.4801 - val_loss: 1.5633
Epoch 27/100
189/189 - 0s - 1ms/step - accuracy: 0.4235 - loss: 1.6885 - val_accuracy: 0.4619 - val_loss: 1.5684
Epoch 28/100
189/189 - 0s - 1ms/step - accuracy: 0.4266 - loss: 1.6908 - val_accuracy: 0.4758 - val_loss: 1.5635
Epoch 29/100
189/189 - 0s - 1ms/step - accuracy: 0.4292 - loss: 1.6584 - val_accuracy: 0.4765 - val_loss: 1.5470
Epoch 30/100
189/189 - 0s - 1ms/step - accuracy: 0.4336 - loss: 1.6726 - val_accuracy: 0.4735 - val_loss: 1.5413
Epoch 31/100
189/189 - 0s - 1ms/step - accuracy: 0.4343 - loss: 1.6584 - val_accuracy: 0.4755 - val_loss: 1.5339
Epoch 32/100
189/189 - 0s - 1ms/step - accuracy: 0.4318 - loss: 1.6580 - val_accuracy: 0.4814 - val_loss: 1.5253
Epoch 33/100
189/189 - 0s - 1ms/step - accuracy: 0.4425 - loss: 1.6469 - val_accuracy: 0.4907 - val_loss: 1.5390
Epoch 34/100
189/189 - 0s - 1ms/step - accuracy: 0.4368 - loss: 1.6565 - val_accuracy: 0.4834 - val_loss: 1.5365
Epoch 35/100
189/189 - 0s - 1ms/step - accuracy: 0.4360 - loss: 1.6557 - val_accuracy: 0.4891 - val_loss: 1.5125
Epoch 36/100
189/189 - 0s - 1ms/step - accuracy: 0.4386 - loss: 1.6449 - val_accuracy: 0.4761 - val_loss: 1.5291
Epoch 37/100
189/189 - 0s - 1ms/step - accuracy: 0.4494 - loss: 1.6235 - val_accuracy: 0.4924 - val_loss: 1.5133
Epoch 38/100
189/189 - 0s - 1ms/step - accuracy: 0.4463 - loss: 1.6294 - val_accuracy: 0.4980 - val_loss: 1.5048
Epoch 39/100
189/189 - 0s - 1ms/step - accuracy: 0.4361 - loss: 1.6412 - val_accuracy: 0.4828 - val_loss: 1.5061
Epoch 40/100
189/189 - 0s - 1ms/step - accuracy: 0.4487 - loss: 1.6275 - val_accuracy: 0.4920 - val_loss: 1.4947

Epoch 41/100
189/189 - 0s - 1ms/step - accuracy: 0.4427 - loss: 1.6314 - val_accuracy: 0.4944 - val_loss: 1.4953
Epoch 42/100
189/189 - 0s - 1ms/step - accuracy: 0.4502 - loss: 1.6142 - val_accuracy: 0.4927 - val_loss: 1.5105
Epoch 43/100
189/189 - 0s - 1ms/step - accuracy: 0.4542 - loss: 1.6136 - val_accuracy: 0.4924 - val_loss: 1.4888
Epoch 44/100
189/189 - 0s - 1ms/step - accuracy: 0.4471 - loss: 1.6194 - val_accuracy: 0.4841 - val_loss: 1.5263
Epoch 45/100
189/189 - 0s - 1ms/step - accuracy: 0.4566 - loss: 1.6068 - val_accuracy: 0.4728 - val_loss: 1.5534
Epoch 46/100
189/189 - 0s - 1ms/step - accuracy: 0.4512 - loss: 1.6081 - val_accuracy: 0.5070 - val_loss: 1.4848
Epoch 47/100
189/189 - 0s - 1ms/step - accuracy: 0.4522 - loss: 1.6023 - val_accuracy: 0.4957 - val_loss: 1.4971
Epoch 48/100
189/189 - 0s - 1ms/step - accuracy: 0.4593 - loss: 1.6012 - val_accuracy: 0.4993 - val_loss: 1.4950
Epoch 49/100
189/189 - 0s - 1ms/step - accuracy: 0.4531 - loss: 1.6054 - val_accuracy: 0.5139 - val_loss: 1.4755
Epoch 50/100
189/189 - 0s - 1ms/step - accuracy: 0.4554 - loss: 1.5912 - val_accuracy: 0.4877 - val_loss: 1.4966
Epoch 51/100
189/189 - 0s - 1ms/step - accuracy: 0.4585 - loss: 1.5880 - val_accuracy: 0.5050 - val_loss: 1.4760
Epoch 52/100
189/189 - 0s - 1ms/step - accuracy: 0.4547 - loss: 1.5855 - val_accuracy: 0.5070 - val_loss: 1.4690
Epoch 53/100
189/189 - 0s - 1ms/step - accuracy: 0.4612 - loss: 1.5841 - val_accuracy: 0.5043 - val_loss: 1.4877
Epoch 54/100
189/189 - 0s - 1ms/step - accuracy: 0.4615 - loss: 1.5845 - val_accuracy: 0.5003 - val_loss: 1.4868
Epoch 55/100
189/189 - 0s - 1ms/step - accuracy: 0.4629 - loss: 1.5819 - val_accuracy: 0.5007 - val_loss: 1.4899
Epoch 56/100
189/189 - 0s - 1ms/step - accuracy: 0.4630 - loss: 1.5769 - val_accuracy: 0.5093 - val_loss: 1.4840
Epoch 57/100
189/189 - 0s - 1ms/step - accuracy: 0.4762 - loss: 1.5583 - val_accuracy: 0.5123 - val_loss: 1.4701
Epoch 58/100
189/189 - 0s - 1ms/step - accuracy: 0.4677 - loss: 1.5684 - val_accuracy: 0.5116 - val_loss: 1.4681
Epoch 59/100
189/189 - 0s - 1ms/step - accuracy: 0.4626 - loss: 1.5699 - val_accuracy: 0.5030 - val_loss: 1.4767
Epoch 60/100
189/189 - 0s - 1ms/step - accuracy: 0.4680 - loss: 1.5755 - val_accuracy: 0.5093 - val_loss: 1.4680

Epoch 61/100
189/189 - 0s - 1ms/step - accuracy: 0.4655 - loss: 1.5663 - val_accuracy: 0.5090 - val_loss: 1.4641
Epoch 62/100
189/189 - 0s - 1ms/step - accuracy: 0.4711 - loss: 1.5545 - val_accuracy: 0.5126 - val_loss: 1.4476
Epoch 63/100
189/189 - 0s - 1ms/step - accuracy: 0.4723 - loss: 1.5444 - val_accuracy: 0.5179 - val_loss: 1.4575
Epoch 64/100
189/189 - 0s - 1ms/step - accuracy: 0.4700 - loss: 1.5622 - val_accuracy: 0.5149 - val_loss: 1.4600
Epoch 65/100
189/189 - 0s - 1ms/step - accuracy: 0.4759 - loss: 1.5468 - val_accuracy: 0.5113 - val_loss: 1.4473
Epoch 66/100
189/189 - 0s - 1ms/step - accuracy: 0.4781 - loss: 1.5514 - val_accuracy: 0.5209 - val_loss: 1.4502
Epoch 67/100
189/189 - 0s - 1ms/step - accuracy: 0.4760 - loss: 1.5448 - val_accuracy: 0.5229 - val_loss: 1.4482
Epoch 68/100
189/189 - 0s - 1ms/step - accuracy: 0.4781 - loss: 1.5385 - val_accuracy: 0.5172 - val_loss: 1.4588
Epoch 69/100
189/189 - 0s - 1ms/step - accuracy: 0.4763 - loss: 1.5428 - val_accuracy: 0.5179 - val_loss: 1.4404
Epoch 70/100
189/189 - 0s - 1ms/step - accuracy: 0.4736 - loss: 1.5356 - val_accuracy: 0.5252 - val_loss: 1.4525
Epoch 71/100
189/189 - 0s - 1ms/step - accuracy: 0.4759 - loss: 1.5434 - val_accuracy: 0.5066 - val_loss: 1.4726
Epoch 72/100
189/189 - 0s - 1ms/step - accuracy: 0.4729 - loss: 1.5415 - val_accuracy: 0.5196 - val_loss: 1.4447
Epoch 73/100
189/189 - 0s - 1ms/step - accuracy: 0.4784 - loss: 1.5425 - val_accuracy: 0.5186 - val_loss: 1.4377
Epoch 74/100
189/189 - 0s - 1ms/step - accuracy: 0.4740 - loss: 1.5432 - val_accuracy: 0.5222 - val_loss: 1.4643
Epoch 75/100
189/189 - 0s - 1ms/step - accuracy: 0.4811 - loss: 1.5332 - val_accuracy: 0.5225 - val_loss: 1.4321
Epoch 76/100
189/189 - 0s - 1ms/step - accuracy: 0.4779 - loss: 1.5390 - val_accuracy: 0.5212 - val_loss: 1.4500
Epoch 77/100
189/189 - 0s - 1ms/step - accuracy: 0.4834 - loss: 1.5300 - val_accuracy: 0.5219 - val_loss: 1.4380
Epoch 78/100
189/189 - 0s - 1ms/step - accuracy: 0.4821 - loss: 1.5199 - val_accuracy: 0.5070 - val_loss: 1.4642
Epoch 79/100
189/189 - 0s - 1ms/step - accuracy: 0.4795 - loss: 1.5294 - val_accuracy: 0.5159 - val_loss: 1.4419
Epoch 80/100
189/189 - 0s - 1ms/step - accuracy: 0.4734 - loss: 1.5224 - val_accuracy: 0.5169 - val_loss: 1.4497

Epoch 81/100
189/189 - 0s - 1ms/step - accuracy: 0.4881 - loss: 1.5033 - val_accuracy: 0.5202 - val_loss: 1.4447
Epoch 82/100
189/189 - 0s - 1ms/step - accuracy: 0.4808 - loss: 1.5129 - val_accuracy: 0.5212 - val_loss: 1.4329
Epoch 83/100
189/189 - 0s - 1ms/step - accuracy: 0.4837 - loss: 1.5262 - val_accuracy: 0.5179 - val_loss: 1.4424
Epoch 84/100
189/189 - 0s - 1ms/step - accuracy: 0.4936 - loss: 1.4941 - val_accuracy: 0.5232 - val_loss: 1.4242
Epoch 85/100
189/189 - 0s - 1ms/step - accuracy: 0.4828 - loss: 1.5225 - val_accuracy: 0.5202 - val_loss: 1.4379
Epoch 86/100
189/189 - 0s - 1ms/step - accuracy: 0.4923 - loss: 1.4993 - val_accuracy: 0.5196 - val_loss: 1.4451
Epoch 87/100
189/189 - 0s - 1ms/step - accuracy: 0.4854 - loss: 1.5171 - val_accuracy: 0.5169 - val_loss: 1.4480
Epoch 88/100
189/189 - 0s - 1ms/step - accuracy: 0.4897 - loss: 1.5069 - val_accuracy: 0.5285 - val_loss: 1.4266
Epoch 89/100
189/189 - 0s - 1ms/step - accuracy: 0.4871 - loss: 1.5023 - val_accuracy: 0.5136 - val_loss: 1.4453
Epoch 90/100
189/189 - 0s - 1ms/step - accuracy: 0.4903 - loss: 1.4866 - val_accuracy: 0.5206 - val_loss: 1.4383
Epoch 91/100
189/189 - 0s - 1ms/step - accuracy: 0.4925 - loss: 1.5015 - val_accuracy: 0.5235 - val_loss: 1.4370
Epoch 92/100
189/189 - 0s - 1ms/step - accuracy: 0.4882 - loss: 1.5082 - val_accuracy: 0.5186 - val_loss: 1.4360
Epoch 93/100
189/189 - 0s - 1ms/step - accuracy: 0.4964 - loss: 1.4898 - val_accuracy: 0.5249 - val_loss: 1.4248
Epoch 94/100
189/189 - 0s - 1ms/step - accuracy: 0.4911 - loss: 1.4975 - val_accuracy: 0.5279 - val_loss: 1.4427
Epoch 95/100
189/189 - 0s - 1ms/step - accuracy: 0.4855 - loss: 1.4981 - val_accuracy: 0.5288 - val_loss: 1.4228
Epoch 96/100
189/189 - 0s - 1ms/step - accuracy: 0.4973 - loss: 1.4901 - val_accuracy: 0.5275 - val_loss: 1.4211
Epoch 97/100
189/189 - 0s - 1ms/step - accuracy: 0.4890 - loss: 1.4920 - val_accuracy: 0.5255 - val_loss: 1.4252
Epoch 98/100
189/189 - 0s - 1ms/step - accuracy: 0.4968 - loss: 1.4908 - val_accuracy: 0.5176 - val_loss: 1.4685
Epoch 99/100
189/189 - 0s - 1ms/step - accuracy: 0.4957 - loss: 1.5008 - val_accuracy: 0.5212 - val_loss: 1.4182
Epoch 100/100
189/189 - 0s - 1ms/step - accuracy: 0.4938 - loss: 1.4967 - val_accuracy: 0.5242 - val_loss: 1.4352

Out[94]: <keras.src.callbacks.history.History at 0x71f6fbfelc90>

```
In [95]: # predict
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
```

118/118 ————— 0s 2ms/step

```
In [96]: # evaluate the model
accuracy = accuracy_score(y_test, y_pred_classes)
precision = precision_score(y_test, y_pred_classes, average='weighted')
recall = recall_score(y_test, y_pred_classes, average='weighted')
f1 = f1_score(y_test, y_pred_classes, average='weighted')
conf_matrix = confusion_matrix(y_test, y_pred_classes)
```

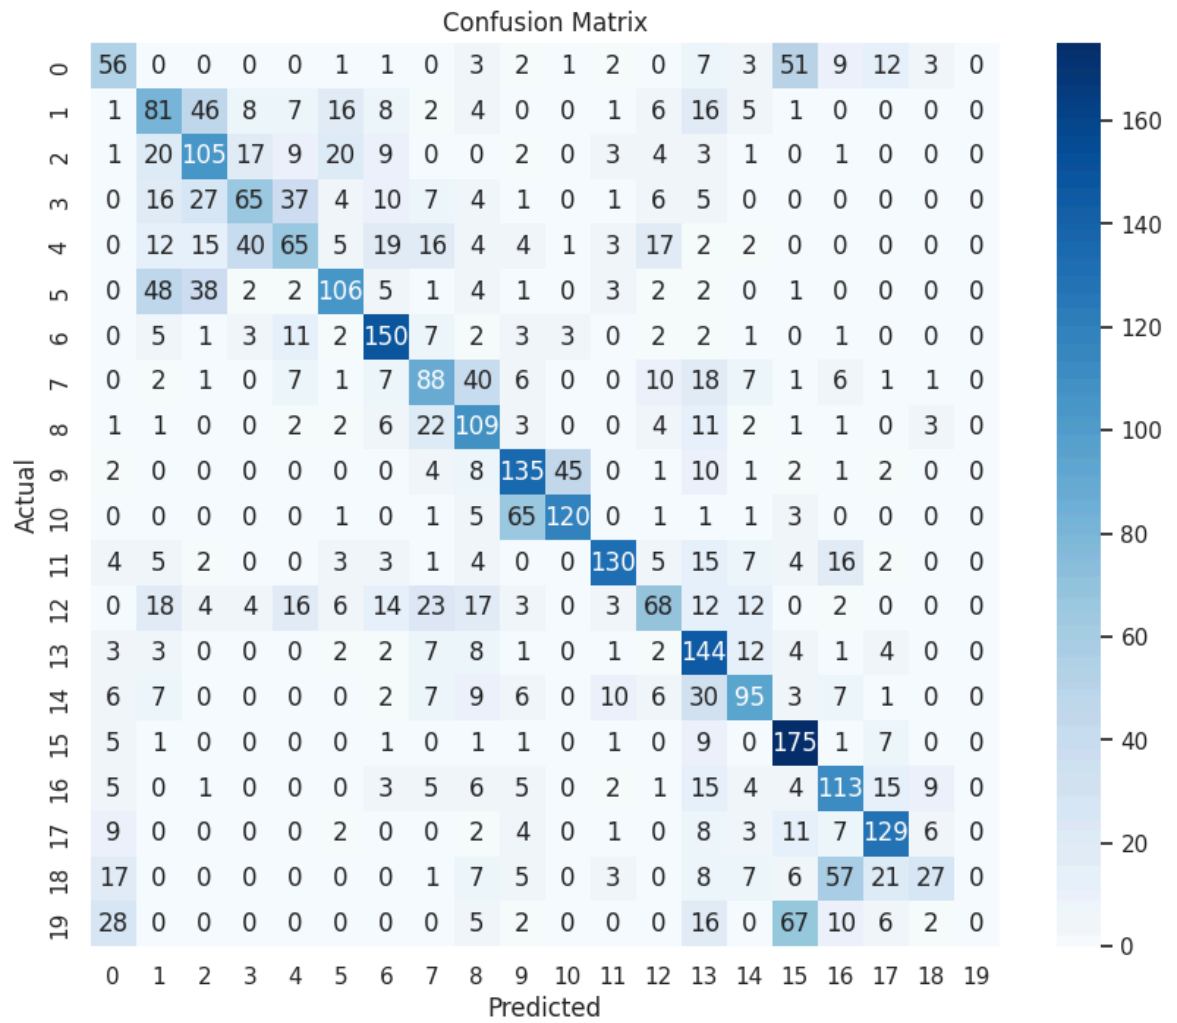
/home/parker/anaconda3/envs/tf_ml/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```
In [101... # Calculate AUC-ROC
y_test_categorical = to_categorical(y_test, num_classes=num_topics)
auc_roc = roc_auc_score(y_test_categorical, y_pred, average='weighted', m
```

```
In [102... print(f'Accuracy: {accuracy}')
print(f'Precision (weighted): {precision}')
print(f'Recall (weighted): {recall}')
print(f'F1 Score (weighted): {f1}')
print(f'Area Under ROC Curve (AUC-ROC): {auc_roc}')
```

Accuracy: 0.520159151193634
Precision (weighted): 0.5091006819709496
Recall (weighted): 0.520159151193634
F1 Score (weighted): 0.5013975361823386
Area Under ROC Curve (AUC-ROC): 0.9420493090049796

```
In [103... plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



In []: